

GRAPHQL



Stanko Krtalic Rusendic

 github.com/Stankec

 [@monorkin](https://twitter.com/monorkin)

Hi! For those who don't know me, my name is Stanko. If you want to see my work or tweet at me here is the needed information.

I've been a Ruby engineer for the better part of the last three years. Lately I've been switching to functional oriented programming and exploring Rust, so perhaps some of you know me from the Lambda Zagreb meetup.



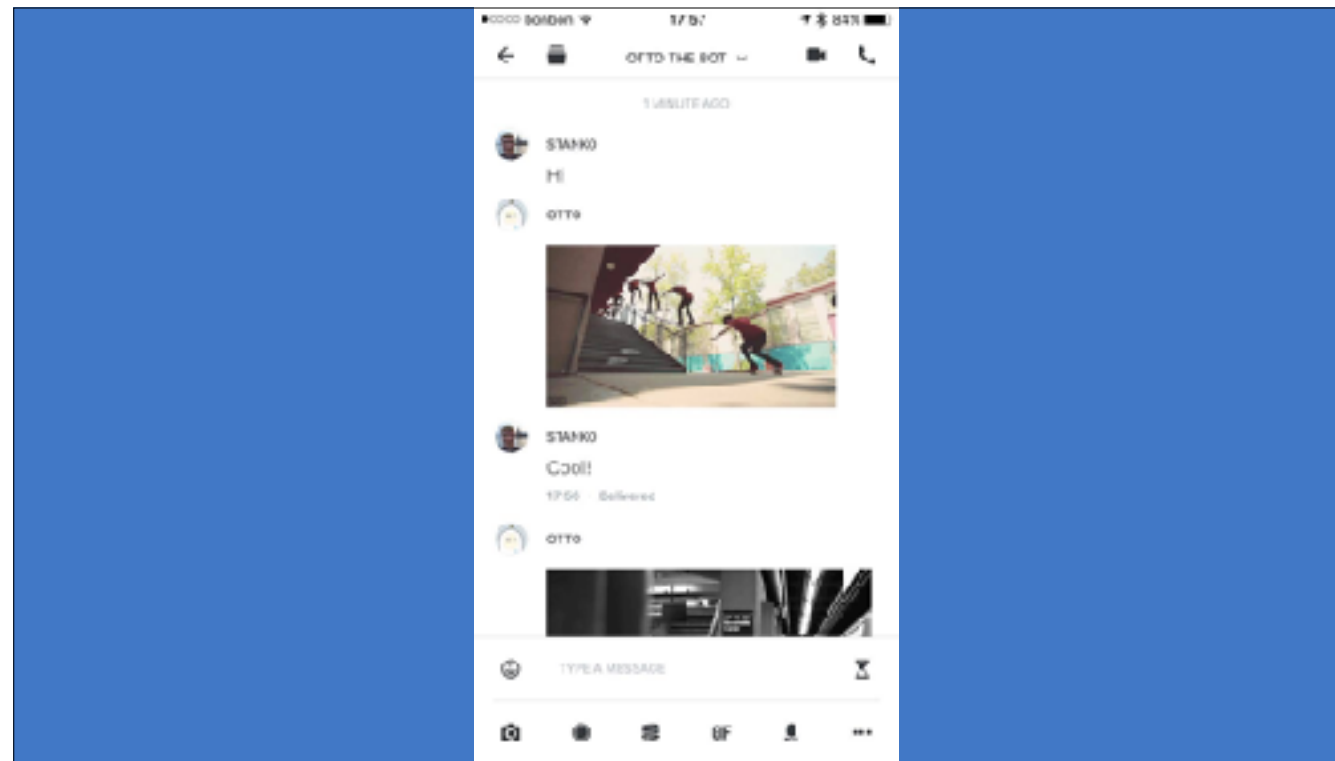
This brings me to my employer, FloatingPoint. If you are interested in working with some of the technologies I've mentioned, Phoenix or Clojure give us a shoutut or nag Nikica and Mihael during the drinkup after the lecture.

01 REST PROBLEMS

So let's begin.
Over the years I've grown tired of REST.

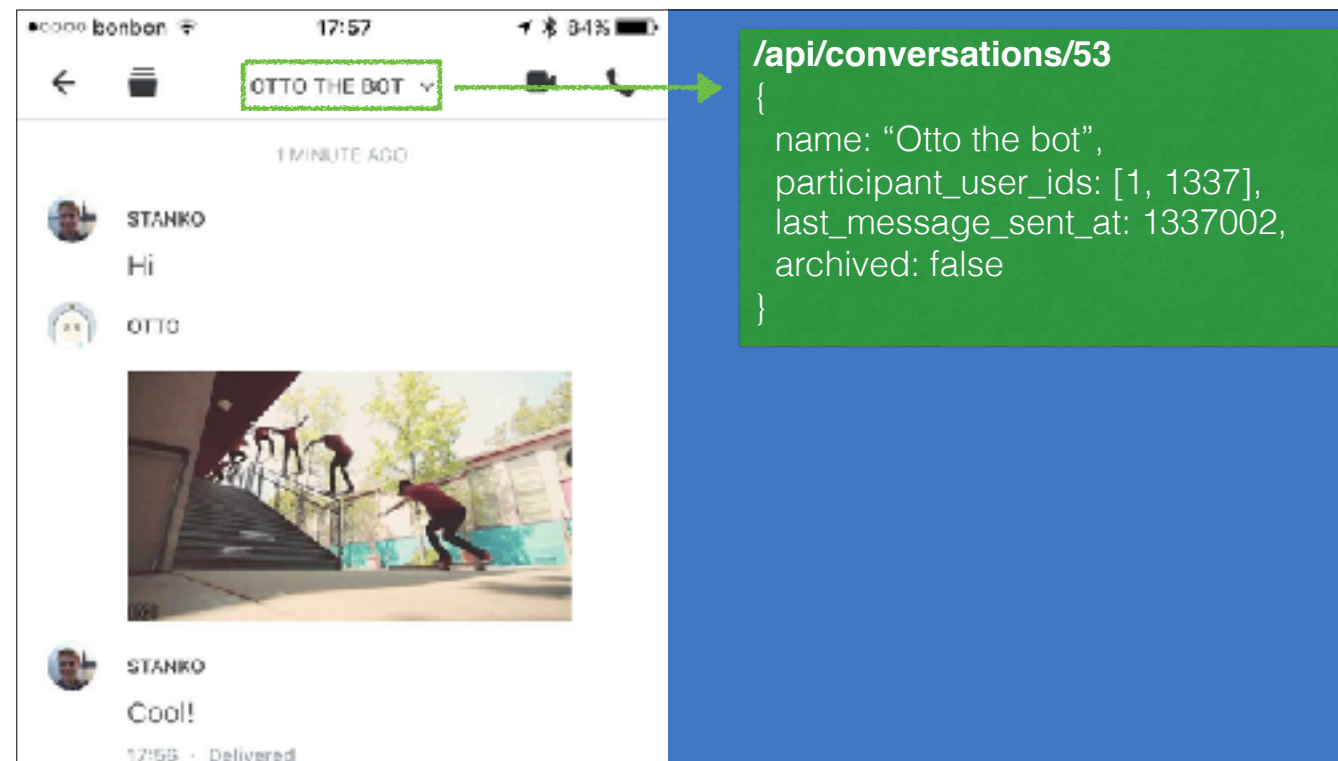
Representational state transfer

I think by now everybody knows this... So REST is a way of organizing your application's resources at different endpoints so that they are easy to explore and retrieve.



Let's go through how you would model a RESTful API for this simple chat application

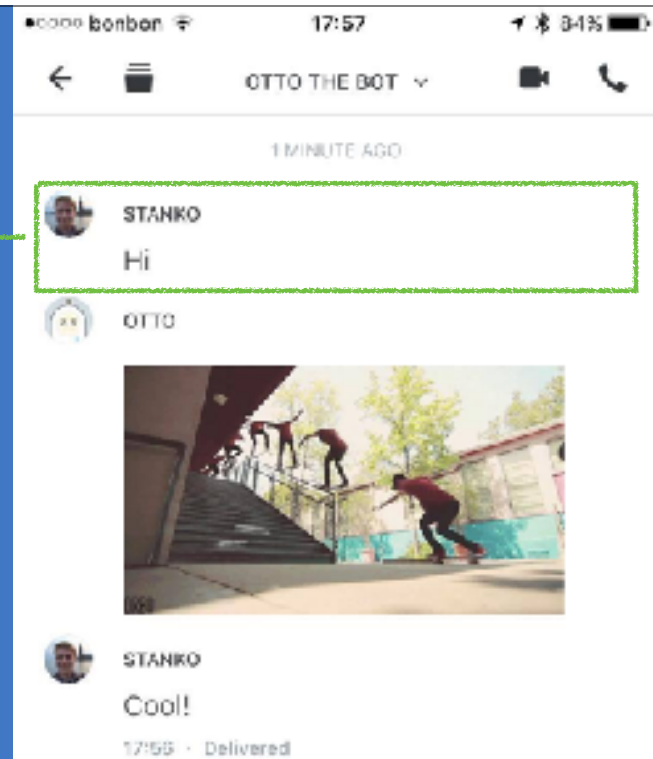
IDEA



You would have a conversations endpoint which serves the name of the conversation, references to it's participants and similar information...

/api/conversations/53/message/1

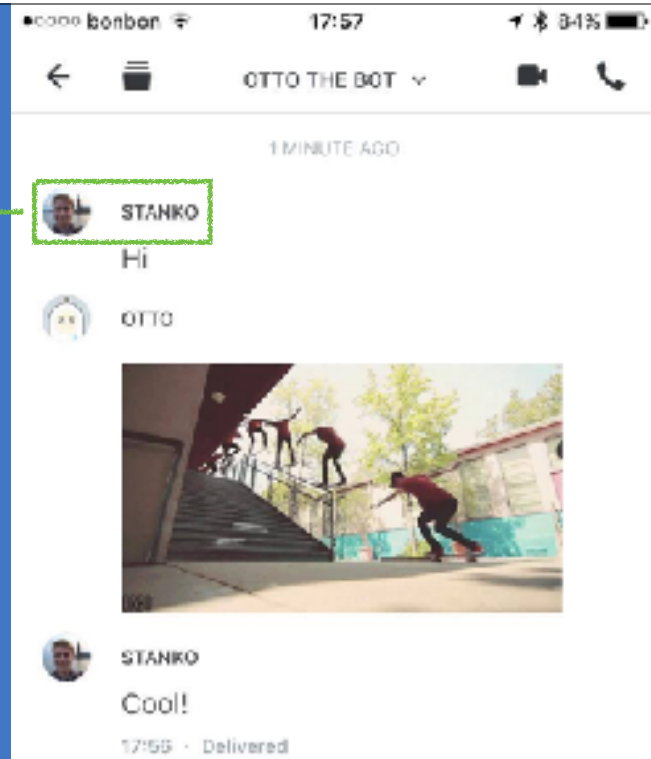
```
{  
  message_type: "plain",  
  body: "Hi",  
  sender_id: "1337",  
  created_at: 1337001,  
  updated_at: 1337001,  
  previous_version_ids: [],  
  status: "delivered",  
  seen_by_participant_ids: [1]  
}
```



Then you would have an endpoint for each individual message...

/api/users/1337

```
{  
  first_name: "Stanko",  
  last_name: "Krtalic Rusendic",  
  avatar_url: "https://...",  
  last_online_at: 1337001,  
  last_signin_ip: 133.10.45.99,  
  sex: "male",  
  timezone: "GMT+1",  
  last_known_location: "16.0E45.0N"  
}
```

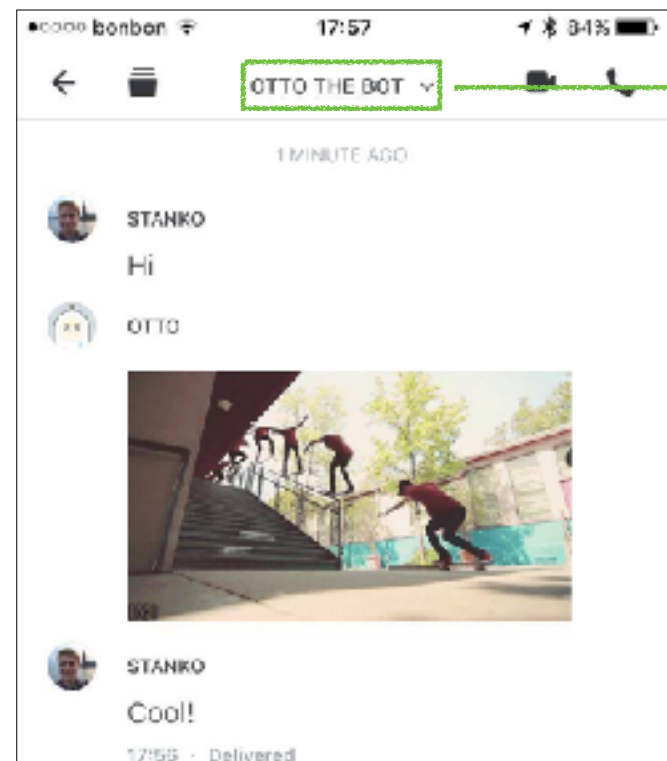


And then you would also create an endpoint to get users.



REALITY

In theory, theory and practice are the same. In practice, not so much...




OTTO THE BOT

1 MINUTE AGO

STANKO
Hi

OTTO



STANKO
Cool!

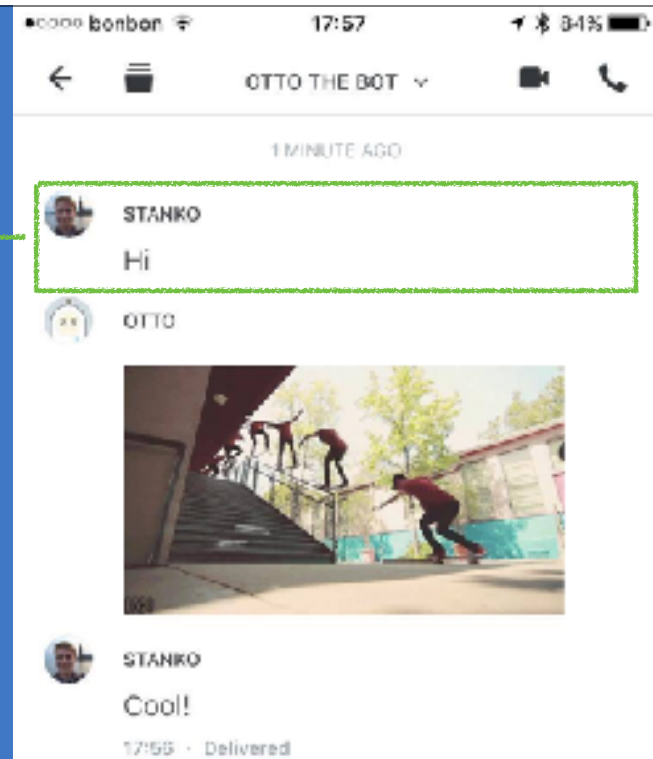
12:55 · Delivered

```
/api/conversations/53  
{  
  name: "Otto the bot",  
  participant_user_ids: [1, 1337],  
  last_message_sent_at: 1337002,  
  last_message_sent_by: {  
    first_name: "Stanko",  
    last_name: "Krtalic Rusendic"  
  }  
  archived: false,  
  last_message: {  
    type: "plain",  
    body: "Cool!",  
    status: "delivered"  
  }  
}
```

Usually the conversations endpoint will serve a lot of information that doesn't actually belong to that resource. E.g. The sender of the latest message with this full data, the content's of the last message, ...

/api/conversations/53/message/1

```
{  
  message_type: "plain",  
  body: "Hi",  
  sender_id: "1337",  
  created_at: 1337001,  
  updated_at: 1337001,  
  previous_version_ids: [],  
  status: "delivered",  
  seen_by_participants: [  
    { first_name: "Otto" }  
  ],  
  sender: {  
    first_name: "Stanko",  
    avatar_url: "https://..."  
  }  
}
```



And messages will contain the sender and a lot of other metainformation...



REST is pointless if your endpoints
respond to a screen in your app

I won't go into as to why this happens, we can blame management for ridiculous time constraints, we can blame frontend developers for being lazy, we can blame backend for giving in too easily, it doesn't matter. It happens and when it happens it breaks REST. If each endpoint returns a lot of data of which some is directly related to the resource and some isn't then you broke REST. This makes development harder because it's much easier to introduce bugs (e.g. infinite resource loops user->articles->author). And in general it makes it harder to separate concerns since one serializer is responsible for a lot of things.

JSON::API HAL

There are some novel efforts to solve this problem.

JSON::API introduces `includes`.

HAL introduced linking.

But in the end both are flawed, one is hard to implement and the other requires a lot of requests for a simple query.

02 DOCUMENTATION

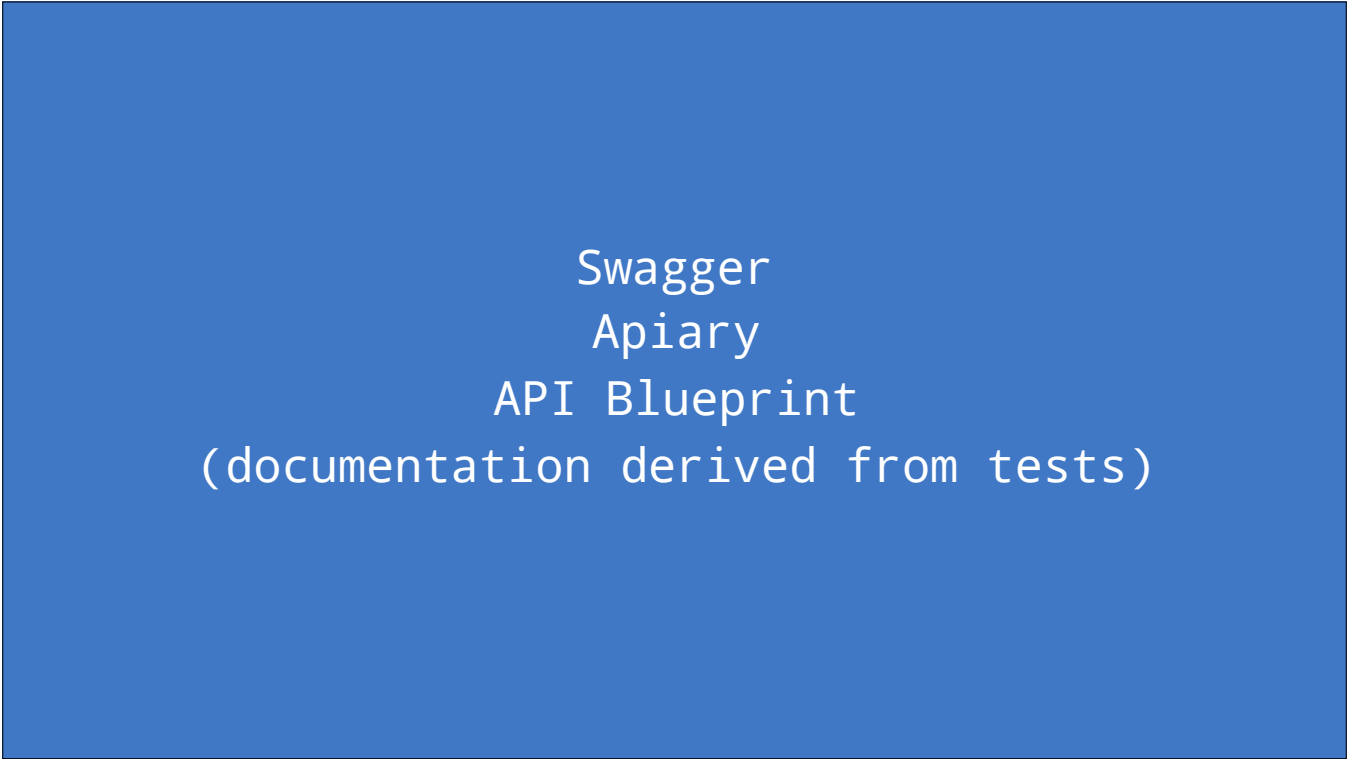
Another problem developers usually face is documentation.



I think everybody was in this situation at least once.

Documentation is boring.

Documentation will always be lacking,
or at least lag behind the
implementation

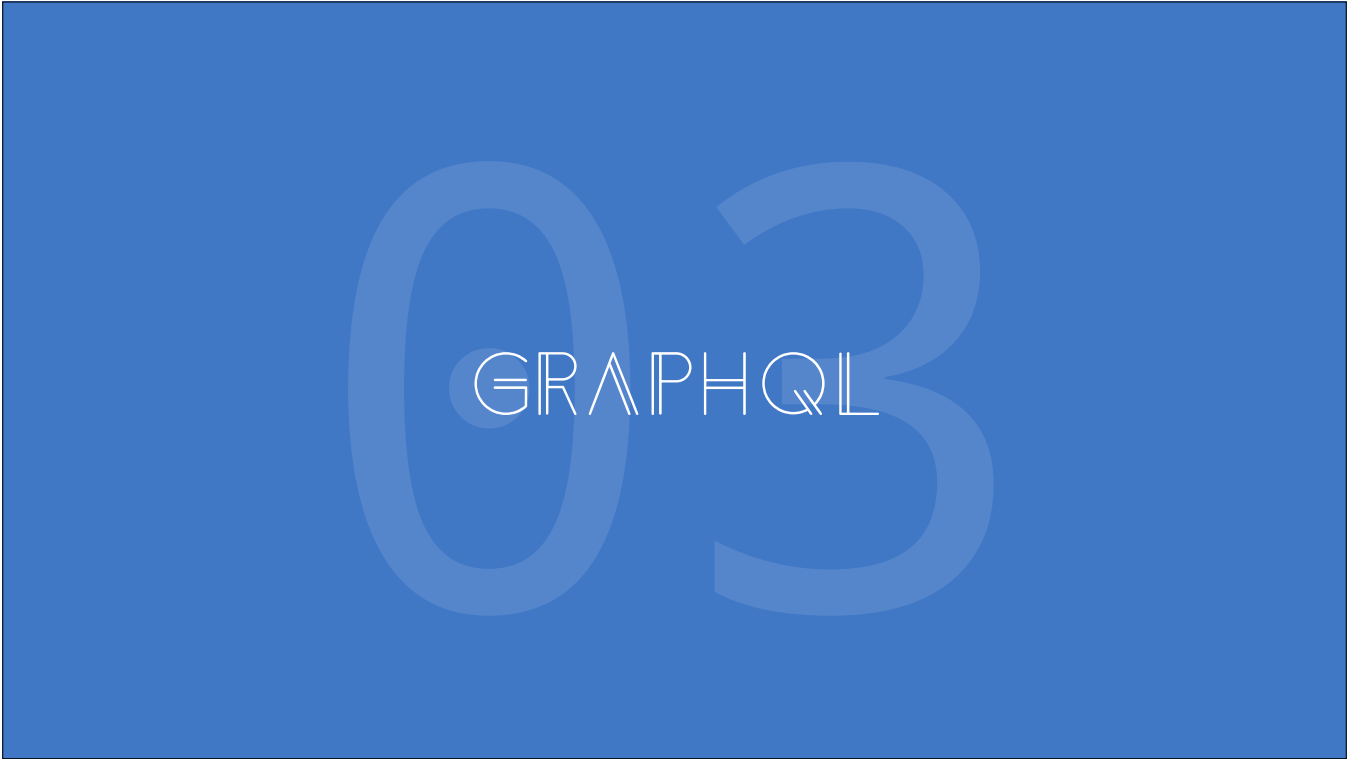


Swagger
Apiary
API Blueprint
(documentation derived from tests)

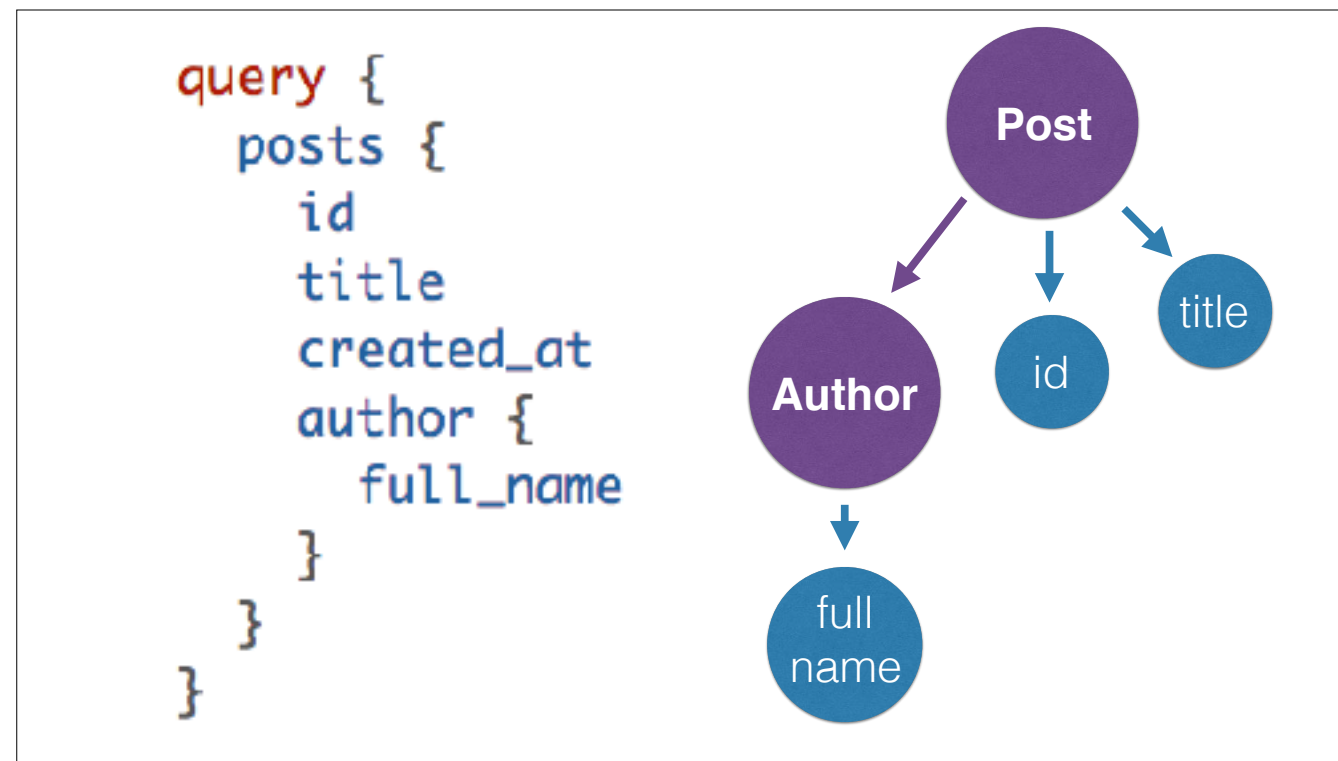
There are a lot of attempts to alleviate this issue.

Require manual labor
Laborious

But none have truly succeeded.



Think of your resources as
endpoints of a graph



As said earlier, each resource or attribute is an endpoint of the graph. Now we just need to specify which endpoints we need.

```
SELECT "posts".id, "posts".title, "posts".body, "posts".created_at  
FROM "posts"  
ORDER BY "posts".created_at DESC
```

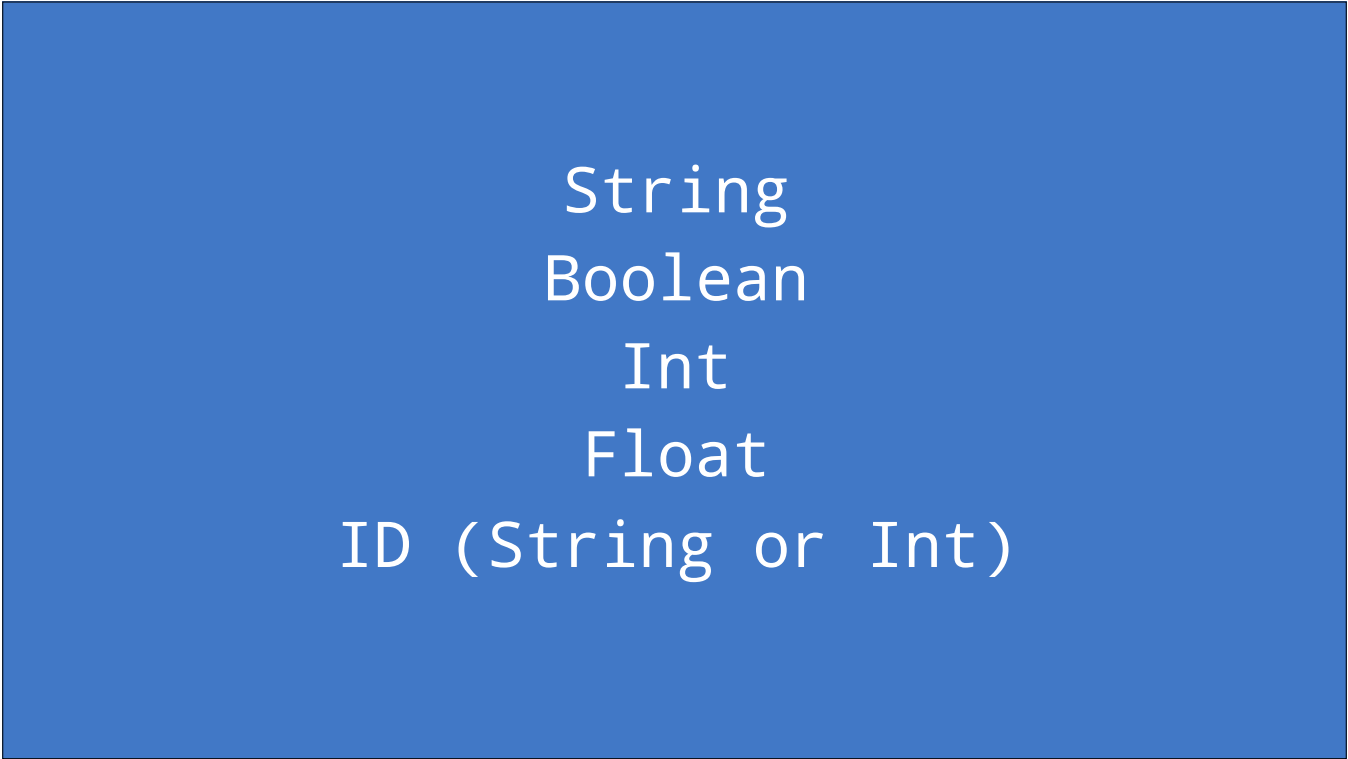
Query languages are especially useful for this purpose. Everybody has had contact with SQL. You basically do exactly what I described a slide ago in SQL, you pick different resources and specify how you want them returned to you.


```
query {  
  posts {  
    id  
    title  
    created_at  
    author {  
      full_name  
    }  
  }  
}
```

```
{  
  "data": {  
    "posts": [  
      {  
        "id": "1",  
        "title": "Lilies of the Field",  
        "created_at": "2017-02-25T18:45:29Z",  
        "author": {  
          "full_name": "Stanko Krstalic Rusendic"  
        }  
      },  
    ]  
  }  
}
```

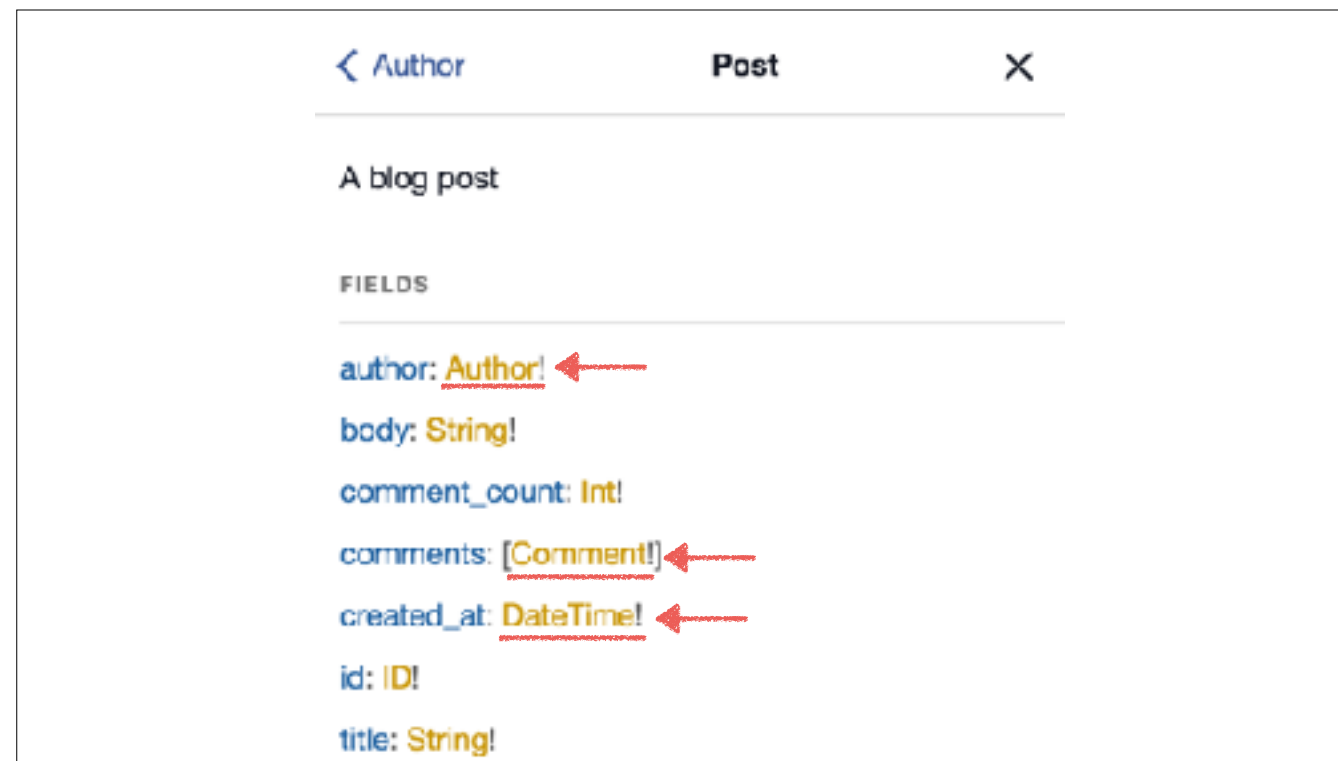
Here is a simple example of a GraphQL query.

I want to demonstrate here that the response will be structured exactly like the query. You can map each attribute from the query to an attribute of the response.



String
Boolean
Int
Float
ID (String or Int)

And before we dig deeper you need to know that GraphQL only supports these basic types.



But you can define custom types.

Queries Mutations

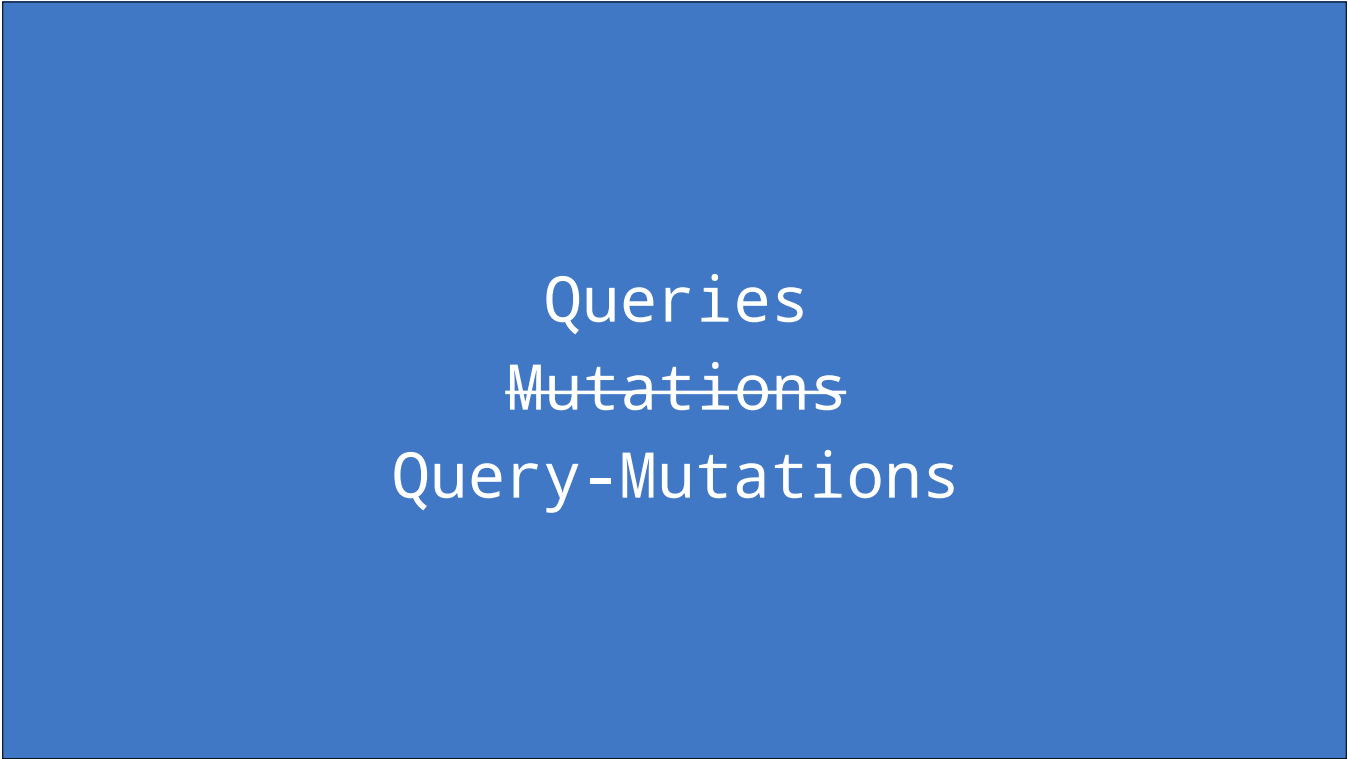
There is also another important thing to note.

In general there are two kinds of methods in programming. Queries and mutations.

```
query {  
  post(id: 1){  
    id  
    title  
    created_at  
    author {  
      full_name  
    }  
    comments {  
      author {  
        first_name  
      }  
      body  
    }  
  }  
}
```

Here is another example of a query

```
{
  "data": {
    "post": {
      "id": "1",
      "title": "Lilies of the Field",
      "created_at": "2017-02-25T18:45:29Z",
      "author": {
        "full_name": "Stanko Krstalic Rusendic"
      },
      "comments": [
        {
          "author": {
            "first_name": "Dario"
          },
          "body": "It was summer.. and it was hot. Rachel was there.. A
lonely grey couch..\"DH LOCK!\" cried Ned, and then the kingdom was his
forever. The End."
        },
        {
          "author": {
            "first_name": "Tomislav"
          },
          "body": "Raspberries? Good. Ladyfingers? Good. BeeF? GOOD!"
        }
      ]
    }
  }
}
```



Queries
~~Mutations~~
Query-Mutations

GraphQL doesn't allow for pure mutations, only for query-mutations. The basic difference being that you have to specify which information you want in return after the query was executed.

```
mutation {  
  createComment(input: {  
    postId: 1  
    authorId: 1  
    body: "This was created using GraphQL 🙌"  
  })  
  {  
    post {  
      comments {  
        body  
      }  
    }  
  }  
}
```

And here is an example of a mutation. Note that you have to select something from the returned information of the mutation. In this example I specify that I want all the contents of all comments on the post I commented.


```
{
  "data": {
    "createComment": {
      "post": {
        "comments": [
          {
            "body": "This was created with GraphQL 🙌"
          }
        ]
      }
    }
  }
}
```

And that is exactly what I get back.

```
query {  
  post_1: post(id: "1") {  
    ...postFields  
  },  
  post_2: post(id: "2") {  
    ...postFields  
  },  
  post_3: post(id: "3") {  
    ...postFields  
  }  
}  
  
fragment postFields on Post {  
  title  
  author {  
    full_name  
  }  
  body  
  comment_count  
}
```

But GraphQL offers many more features. E.g. these are called fragments. Let's say you want the first three posts with the exact same data returned. You would specify a fragment with the data you wish to have returned and simply include it in the query.

```

{
  "data": {
    "post_1": { ➡ },
    "post_2": { ➡ },
    "post_3": {
      "title": "I Sing the Body Electric",
      "author": {
        "full_name": "Stanko Krtalic Rusendic"
      },
      "body": "If we override the driver, we can get to the GB
application through the multi-byte XML alarm!\nWe need to index the 1080p
RAM interface!\nTry to transmit the AI bus, maybe it will connect the
digital interface!\nYou can't connect the interface without connecting
the solid state SQL driver!\nUse the optical EXE bus, then you can
compress the haptic alarm!\nThe PCI interface is down, input the open-
source capacitor so we can back up the SMTP feed!\nOverriding the pixel
won't do anything, we need to synthesize the virtual usb card!\nI'll
connect the 1080p RSS driver, that should bandwidth the GB array!",
      "comment_count": 5
    }
  }
}

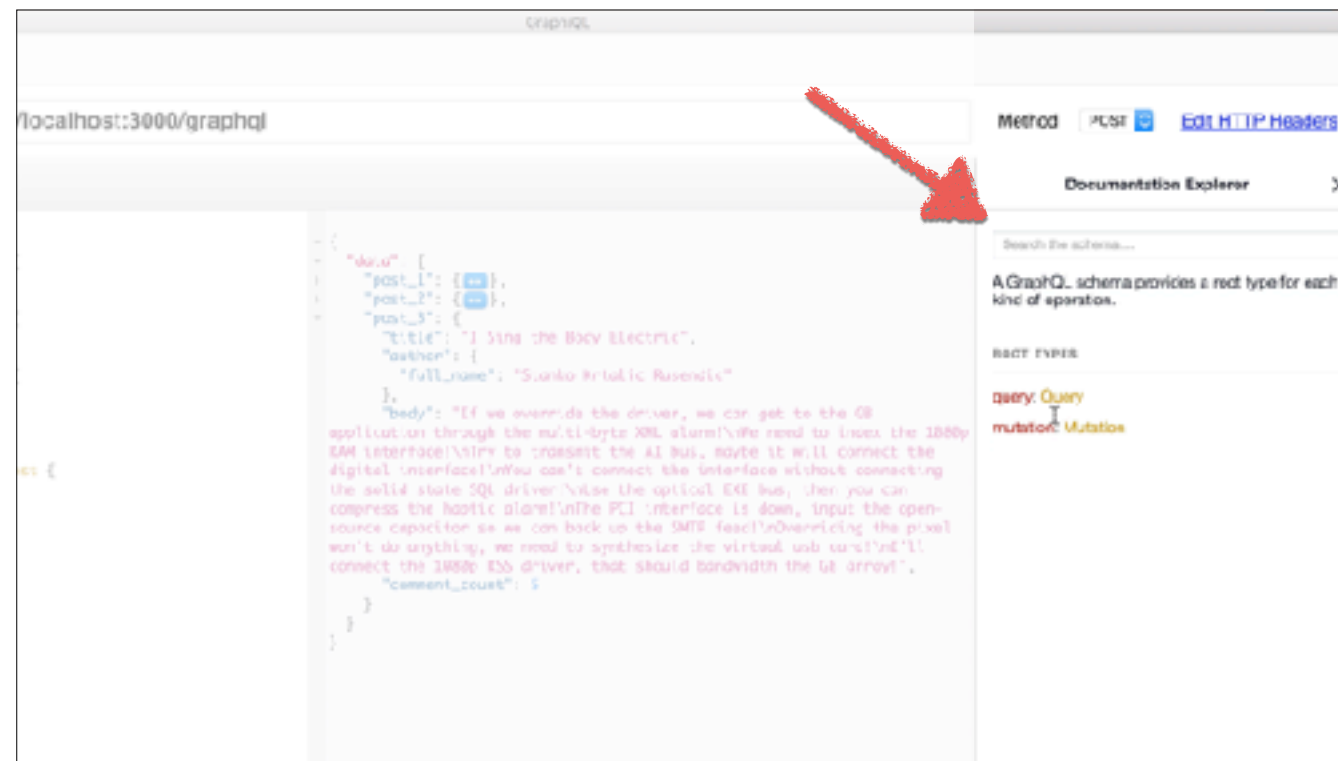
```

This is the result of the query. The blu arrows are just collapsed data. All responses are by structure identical to the one displayed.

```
{
  search(text: "an") {
    __typename
    ... on Human {
      name
    }
    ... on Droid {
      name
    }
    ... on Starship {
      name
    }
  }
}
```

```
{
  "data": {
    "search": [
      {
        "__typename": "Human",
        "name": "Han Solo"
      },
      {
        "__typename": "Human",
        "name": "Leia Organa"
      },
      {
        "__typename": "Starship",
        "name": "TIE Advanced x1"
      }
    ]
  }
}
```

There is also support for polymorphic selections. If some types have different attributes that you want selected but are returned in the same query.



And perhaps the best feature! Autogenerated documentation!

04 GRAPHQL IN RUBY

Ok, but how do we do it in ruby?

```
21 gem 'kaminari'  
22 gem 'trix'  
23 gem 'graphql'  
24 gem 'graphql-batch'  
25 gem 'shrine', '~> 2.2'
```

You add the gem

```
1 Rails.application.routes.draw do
2   root to: 'posts#index'
3   resources :comments
4   resources :posts
5   resources :authors
6
7   # You only need this for GraphQL
8   resource :graphql, only: [:create], controller: 'graphql'
9 end
```

Add the route


```
1 class GraphQLController < ApplicationController
2   skip_before_action :verify_authenticity_token
3
4   def create
5     query_string = params[:graphql][:query]
6     variables = params[:graphql][:variables].try(:to_unsafe_hash)
7     context = {
8       # current_user: current_user
9     }
10
11     graphql_response = GraphQLSchema.execute(
12       query_string,
13       variables: variables,
14       context: context
15     )
16
17     render json: graphql_response
18   end
19 end
```

Create a simple controller.

Note that here you can specify a context, e.g. the logged in user, or some headers that have been passed, or whatever you would like. This will later be available globally to all queries, mutations, resolvers, and whatnot.

```
1 GraphQLSchema = GraphQL::Schema.define do
2   query Graph::Types::QueryType
3   mutation Graph::Types::MutationType
4 end
```

Implement a schema.

A schema is the entrypoint of your GraphQL resolver. It specifies what your app can and can't do.

Each schema has to have at least a query or a mutation type.

Those define what queries and mutations are available.

```

1 module Graph
2   module Types
3     QueryType = GraphQL::ObjectType.define do
4       name 'Query'
5       description 'The query root of this schema'
6
7       # Individual getters
8       field :post, PostType,
9         field: Queries::BasicGetter.by_id(type: PostType, model: Post)
10      field :author, AuthorType,
11        field: Queries::BasicGetter.by_id(type: AuthorType, model: Author)
12      field :comment, CommentType,
13        field: Queries::BasicGetter.by_id(type: CommentType, model: Comment)
14
15      # Batch getters
16      field :authors, types[!AuthorType],
17        field: Queries::BasicGetter.batch(
18          type: types[!AuthorType],
19          model: Author
20        )
21      field :comments, types[!CommentType], field: Queries::CommentsQuery
22      field :posts, types[!PostType],
23        field: Queries::BasicGetter.batch(
24          type: types[!PostType],
25          model: Post
26        )
27    end
28  end
29 end

```

An example of a query type.
It specifies what kind of queries exist.

```
1 module Graph
2   module Types
3     MutationType = GraphQL::ObjectType.define do
4       name 'Mutation'
5       description 'The mutation root of this schema'
6
7       field :createComment, field: Mutations::CreateCommentMutation.field
8     end
9   end
10 end
```

And a mutations type which specifies what kind of mutations exist.

```
1 module Graph
2   module Types
3     PostType = GraphQL::ObjectType.define do
4       name 'Post'
5       description 'A blog post'
6       field :id, !types.ID
7       field :title, !types.String
8       field :body, !types.String
9       field :author, !AuthorType
10      field :comments, types[!CommentType]
11      field :comment_count, !types.Int
12      field :created_at, !Scalars::DateTimeScalar
13    end
14  end
15 end
```

And then you need to specify a type for each resource you have.

You can think of types as serializers. They explain which attributes from the object should be exposed.

You can also deprecate fields here, or define inline resolvers.

Organization

Touching the subject of resolvers, let's talk about code organisation.

```
1 module Graph
2   module Mutations
3     CreateUserMutation = GraphQL::Relay::Mutation.define do
4       name 'CreateUser'
5
6       input_field :email, types.String
7       input_field :password, types.String
8       input_field :fullName, types.String
9       input_field :firstName, types.String
10      input_field :lastName, types.String
11      input_field :phone, types.String
12      input_field :receivePromotionalEmails, types.Boolean
13
14      return_field :user, Types::UserType
15      return_field :token, types.String
16      return_field :errors, Scalars::JsonScalar
17      return_field :error_messages, types[types.String]
18
19      resolve Resolvers::User::Creator
20    end
21  end
22 end
```

Instead of passing Lambdas as specified in the documentation you can also pass objects or classes.

```
1 module Graph
2   module Resolvers
3     module User
4       class Creator
5         attr_reader :object
6         attr_reader :inputs
7         attr_reader :ctx
8
9         def self.call(object, inputs, ctx)
10           new(object, inputs, ctx).call
11         end
12
13         def initialize(object, inputs, ctx)
14           @object = object
15           @inputs = inputs
16           @ctx = ctx
17         end
18       end
19     end
20   end
21 end
```

The passed class has to have a call method and has to accept three parameters - the object, the inputs and a context.


```
▼ [x]app/  
  ▶ assets/  
  ▶ channels/  
  ▶ controllers/  
  ▶ data_objects/  
  ▼ [x]graphql/  
    ▼ [x]graph/  
      ▶ lazy_executors/  
      ▶ mutations/  
      ▶ queries/  
      ▶ [x]resolvers/  
      ▶ scalars/  
      ▶ types/  
      graph_ql_schema.rb
```

I lifted the naming resolvers form the official JS GraphQL implementation.

Optimizing

This is an often question I get. How do I optimize this, or how do I resolver N+1 queries.

```
1 GraphQLSchema = GraphQL::Schema.define do
2   query Graph::Types::QueryType
3   mutation Graph::Types::MutationType
4
5   lazy_resolve(Graph::LazyExecutors::Single, :resolve)
6   lazy_resolve(Graph::LazyExecutors::Batch, :resolve)
7 end
```

The schema allows for definition of lazy resolvers. Basically you tell it which method to call on an object of the specified class if it encounters it during serialization.

```

1 module Graph
2   module Types
3     ProductType = GraphQL::ObjectType.define do
4       name 'Product'
5       description 'A product'
6       field :id, types.ID
7       field :name, types.String
8       field :sku, types.String
9       field :available_on, Scalars::DateTimeScalar
10      field :description, types.String
11      field :meta_description, types.String
12      field :meta_keywords, types.String
13      field :meta_title, types.String
14      field :discontinue_on, Scalars::DateTimeScalar
15      field :prices, types[PriceType]
16      field :variants, types[VariantType]
17      field :variant_images, types[ImageType]
18      field :taxonomies, types[Types::BrandType]
19      field :created_at, Scalars::DateTimeScalar
20      field :updated_at, Scalars::DateTimeScalar
21      field :price, types.Float
22      field :currency, types.String
23      field :recommendations, types[Types::ProductType] do
24        resolve ->(obj, args, ctx) do
25          LazyExecutors::Batch
26            .new(Spree::Product, ctx, obj.recommended_product_ids)
27        end
28      end
29    end
30  end
31 end

```

So here is both an example of an inline resolver and a lazy executor.

```

 9  def initialize(model_class, query_ctx, id)
10    @model_class = model_class
11    @query_ctx = query_ctx
12    @id = id
13
14    @lazy_state = query_ctx[:lazy_find_product] ||= {
15      pending_ids: Set.new,
16      loaded_ids: {},
17    }
18
19    lazy_state[:pending_ids] << id
20  end
21
22  def resolve
23    loaded_record = lazy_state[:loaded_ids][id]
24
25    return loaded_record if loaded_record
26
27    pending_ids = lazy_state[:pending_ids].to_a
28    records = scope.where(id: pending_ids)
29    records.each { |record| lazy_state[:loaded_ids][record.id] = record }
30    lazy_state[:pending_ids].clear
31    lazy_state[:loaded_ids][id]
32  end

```

This is the code of the executor. When it's initialized it remembers which id it has to load as well as which ids it has already loaded. When resolve gets called it either executes a query to fetch all pending ids at once (if the object wasn't found) or it returns the found object.

```

19   def self.batch(model:, type:, scope: nil)
20     return_type = type
21     GraphQL::Field.define do
22       type(return_type)
23       argument(
24         :page, types.Int,
25         "Index of the page, by default only 10 elements are shown per page"
26       )
27       description("Return all #{model.name.pluralize}")
28       resolve _->(obj, args, ctx) {
29         new_scope = scope && scope.respond_to?(:call) &&
30           _scope.call(obj, args, ctx)
31         scope = new_scope || scope
32         relation = scope || model.all
33         GraphQL::Loader.call(ctx, relation)
34       }
35     end
36   end

```

Or you can roll your own. The previous example only solves problems if you can get the id's without making a query. But if you have a has_many relationship that might not be possible.

I've created a utility that preloads all data from your GraphQL query to your SQL query. It uses the Rails includes method in the background. This it's publically available yet, but will be by next Monday.

Questions?

Any questions?